

COMP-40730 HPC

REPORT FOR ASSIGNMENT 3

Author: Paula Dwan
Due date: 13-July-2014
Lecturer: Alexey Lastovetsky
Subject: COMP-40730 High Performance Computing
College: University College Dublin

CONTENTS

Contents.....	2
Exercise.....	3
Overview of Computations Obtained and How.....	4
Assignment Execution.....	5
Running A3-omp-1D : standalone.....	5
Log Files Obtained.....	5
GNUplot Execution.....	7
Summary Results :.....	9
GNUplot Graphs.....	9
Conclusions.....	10
Appendix I – Validate Results.....	12
Appendix II – Acknowledgements.....	13

EXERCISE

Write a parallel OpenMP program computing the norm of the product of two $n \times n$ dense matrices on a p -processor SMP so that

- p threads are involved in the parallel computations.
- The 1-dimensional parallel algorithm of matrix multiplication is employed:
 - one of matrices is partitioned in one dimension into p equal slices
 - there is one-to-one mapping between the partitions and threads
 - each thread is responsible for computation of the corresponding slice of the resulting matrix
- Computation of the norm of the resulting matrix employs the mutex synchronization mechanism.

You can use BLAS or ATLAS for local computations.

Experiment with the program and build:

- The dependence of the execution time of the program on the matrix size n .
- The speedup over a serial counterpart of the program.

Explain the results.

Variants of the assignment:

1. Granularity of the program:
 - (a) Two successive steps:
 - i. Parallel matrix multiplication
 - ii. Parallel computation of the norm of the resulting matrix
 - (b) One-step algorithm. No intermediate resulting matrix.

2. Partitioning scheme:
 - (a) Left matrix is horizontally partitioned
 - (b) Right matrix is vertically partitioned

3. Matrix norm to be computed:
 - (a) The maximum absolute column sum norm (aka one-norm):

- (b) The maximum absolute row sum norm (aka infinity-norm):

$$\|A\|_{\infty} = \max_{0 \leq i < n} \sum_{j=0}^{n-1} |a_{ij}|$$

OVERVIEW OF COMPUTATIONS OBTAINED AND HOW

Assignment 3 basically involved (for me) writing one program which utilized open MP when calculating manually and BLAS when calculating otherwise. :

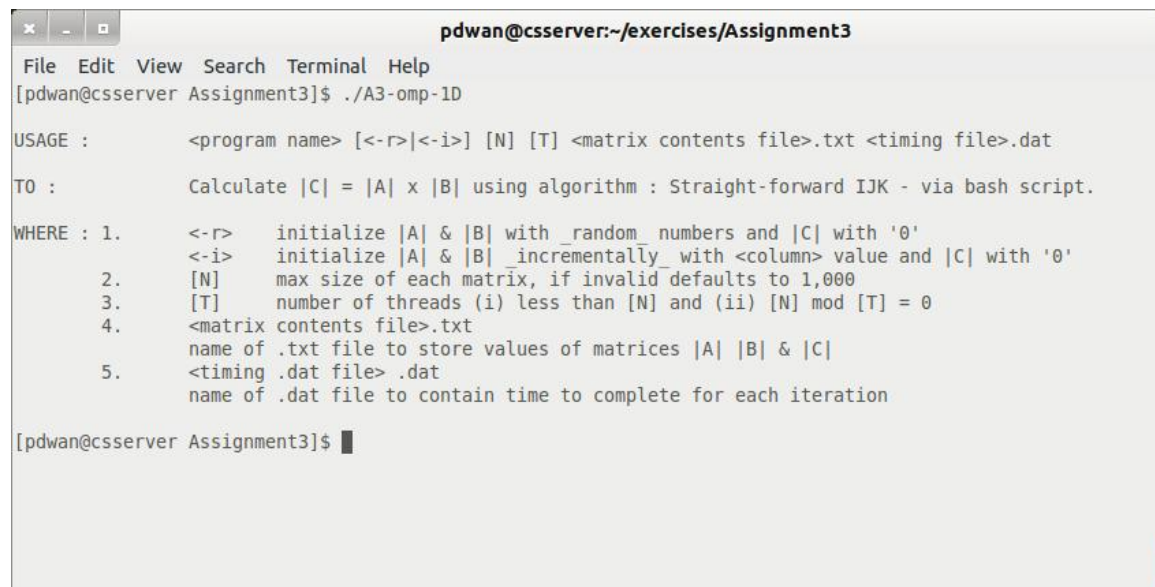
A3-omp-1D.c

Algorithm
using additional
variable

```
for (i=0; i<rows; i++)
{
    for (j=0; j<cols; j++)
    {
        double sum = 0.0;
        for (k=0; k<rows; k++)
        {
            sum+= (A [i][k]) * (B [k][j]);
        }
        C [i][j] = sum;
    }
}
```

|C| matrix was calculated using the algorithm as covered in the first assignment. The same matrix computation was implemented using cblas. Time taken to calculate |C| and the infinity norm was noted and graphed.

This was completed using **A3-omp-1D.c** :



```
pdwan@cssserver:~/exercises/Assignment3
File Edit View Search Terminal Help
[pdwan@cssserver Assignment3]$ ./A3-omp-1D

USAGE :      <-r>|<-i> [N] [T] <matrix contents file>.txt <timing file>.dat

TO :         Calculate |C| = |A| x |B| using algorithm : Straight-forward IJK - via bash script.

WHERE : 1.    <-r>   initialize |A| & |B| with _random_ numbers and |C| with '0'
           <-i>   initialize |A| & |B| _incrementally_ with <column> value and |C| with '0'
           [N]    max size of each matrix, if invalid defaults to 1,000
           [T]    number of threads (i) less than [N] and (ii) [N] mod [T] = 0
           <matrix contents file>.txt
           name of .txt file to store values of matrices |A| |B| & |C|
           5.    <timing .dat file> .dat
           name of .dat file to contain time to complete for each iteration

[pdwan@cssserver Assignment3]$
```

This is compiled using *gcc* and *openmp*:

```
[pdwan@cssserver Assignment3]$ gcc -I/home/cs/khasanov/libs/CBLAS/src A3-omp-1D.c -o A3-omp-1D
/home/cs/khasanov/libs/cblas_LINUX.a /usr/lib/libblas.a -lgfortran -fopenmp
```

and executed on a once off using (for example)

```
[pdwan@cssserver Assignment3]$ ./A3-omp-1D -r 10 2 matrix2.txt timings2.dat
```

Where **-r** indicates random number generation (for 1 to 10) and **-i** indicates number of column + 1.

ASSIGNMENT EXECUTION

Each program was executed multiple times using the script `./runAssignment3.sh`. This has multiple options and the syntax and usage follows:

```
UCD - Assignments
File Edit View Search Terminal Help
(master)~/gitrepos/UCD/COMP40730.HPC/Assignment3$ ./runAssignment3.sh

USAGE : ./runAssignment3.sh \
        -d2|--cblas -r|--random -i|--increment -m|--matrix <n> -t|--thread <t> -v|--values -?-h|--help

TO : Calculate |C| = |A| x |B| and then infinity norm using Open MP

LOGS : Created in current dir and moved to <logDir> :
        <file>.txt : matrix values for matrices |A| |B| & |C|,
        <file>.dat : timing data for each computation
        <file>.log : summary of stdout.

WHERE : -d2|--cblas Compile .c source files using dgemm cblas & openmp
        -r|--random Initialize |A| & |B| with random numbers and |C| with '0'
        -i|--increment Initialize |A| & |B| incrementally with <column> value and |C| with '0'
        '-i|--increment' & '-r|--random' are mutually exclusive

        -m|--matrix <n> Matrix dimension, if odd number +1 added or if invalid set to [ 1000 ], thread count set to [ 20 ]
        -t|--thread <t> number of threads, if invalid set to [ 20 ] and matrix size set to [ 1000 ]
        -v|--values Use predefined range of valid values for <nx> and <nb> as follows :
        <matrixArray> : { 50, 50, 50, 100, 100, 100, 500, 500, 500, 1000, 1000, 1000 }
        <threadArray> : { 10, 10, 10, 10, 10, 20, 20, 20, 20, 20, 20 }
        '-m|--matrix <n>' | '-t|--thread <t>' and '-v|--values' mutually exclusive.

        -?-h|--help usage

(master)~/gitrepos/UCD/COMP40730.HPC/Assignment3$
```

Execute this script in the home directory of Assignment 3.

Note: Please retain the overall directory structure when unzipping.

Note that the script `./runAssignment3.sh` allows two types of implementation

- Multiple iteration : use the switch `<-v|--values>`, when a predefined range applies for [N] : matrix size and [T] : number of threads applicable.
- Single iteration : use the switch `<-m|--matrix> [N]` where the user specifies the values for [N] : matrix size and [T] : number of threads applicable.

RUNNING A3-OMP-1D : STANDALONE

The compiled .c program may also be run standalone. Usage follows:

```
pdwan@cserver:~/exercises/Assignment3
File Edit View Search Terminal Help
[pdwan@cserver Assignment3]$ ./A3-omp-1D

ERROR : <number of arguments> : 1, is invalid, less than <default> : 6.

USAGE : <program name> [<-r>|<-i>] [N] [T] <matrix contents file>.txt <timing file>.dat

TO : Calculate |C| = |A| x |B| using Open MP and also calculate infinity norm of |C|

WHERE : 1. <-r> initialize |A| & |B| with _random_ numbers and |C| with '0'
        <-i> initialize |A| & |B| incrementally with <column> value and |C| with '0'
        2. [N] max size of each matrix, if invalid defaults to 1,000
        3. [T] number of threads (i) less than [N] and (ii) [N] mod [T] = 0
        4. <matrix contents file>.txt
        name of .txt file to store values of matrices |A| |B| & |C|
        5. <timing .dat file> .dat
        name of timing data file to containing calculation time for each iteration

[pdwan@cserver Assignment3]$
```

LOG FILES OBTAINED

Data text files suitable containing the values of the computation used for matrices |A| and |B| and the results stored in |C| are saved in the appropriate log files. File naming convention via the script is :

<data log file name>	pdwan-<time>-values-A3-omp-1D->-<iteration>.txt
example:	pdwan-20140708.015835-values-A3-omp-1D-8.txt

A summary file containing processing time for each computation (manual and BLAS) for is also saved. This is in a format suitable for us with GNUpot.

example: pdwan-20140708.015835-timing-A3-omp-1D-8.dat

If ***./A3-omp-1D*** is used without the script then files may be named whatever the user wishes and no .log file applies.

```

pdwan@csserver:~/exercises/Assignment3
File Edit View Search Terminal Help
# -----
#
# Program :      A3-omp-1D
# where :      .dat contains timing data & .txt contains matrix values
# -----
# |Matrix|      |Threads|      Time/manual      Inf Norm/manual      Time/dgemm      Inf Norm/dgemm
# -----
6      3      0.001521s      441      0.001662s      441
6      3      0.001246s      1396      0.001480s      1396

```

```

pdwan@csserver:~/exercises/Assignment3
File Edit View Search Terminal Help
#
# RUNNING :      ./A3-omp-1D -r 4 2  t7.dat
# Program :      A3-omp-1D
# where :        .dat contains timing data & .txt contains matrix values
#
# Summary of values added to each matrix - retained for later reference and validation
#
# -----
# Initialize results <4> x <4> |A| ...
4      7      8      6
4      6      7      3
10     2      3      8
1      10     4      7

# Initialize results <4> x <4> |B| ...
1      7      3      7
2      9      8      10
3      1      3      4
8      6      10     3

# Initialize results <4> x <4> |C| ...
0      0      0      0
0      0      0      0
0      0      0      0
0      0      0      0

# RESULTS : calculation where number of threads are : 2

# |C| : <4> x <4> matrix computed values : MANUAL ...
90     135     152     148
61     107     111     125
87     139     135     126
89     143     165     144

```

I wished to keep each .c program as clean as possible and so all production setup was completed in the script for each assignment. Thus file creation and validation for each iteration was completed before the .c program was even called. Simple validation of the arguments passed to each .c program is also completed.

I also spot-checked the results as practical. Results obtained are detailed in [Appendix I – Validate Results](#).

GNUplot EXECUTION

I followed the same convention for each .dat file as produced, an example follows :

Sample .dat file	<pre> pdwan@csserver:~/exercises/Assignment3 File Edit View Search Terminal Help ----- # # Program : A3-omp-1D # where : .dat contains timing data & .txt contains matrix values # # Matrix Threads Time/manual Inf Norm/manual Time/dgemm Inf Norm/dgemm # 6 3 0.001521s 441 0.001662s 441 6 3 0.001246s 1396 0.001480s 1396 </pre>					
	<p style="text-align: right;">1,1 Top</p>					

Each was then presented in graphical format using GNUplot, comparing times taken for manual and for BLAS/ATLAS computations. A generic GNUplot program was written to output the data to the screen.

GNUplot program execution	<pre> # To execute, launch GNUplot and run : # gnuplot> load <filename.gp> # making sure that the data file name used is updated if needed. # Paula Dwan : Assignment 3 : A3-plotgraph-matrix.gp reset set xtic auto set ytic auto set size 1,1 set grid set key outside # set title 'Comparison : Matrix Size v time taken' set ylabel 'Time taken / s' set xlabel 'Matrix size' set origin 0,0 plot 'logDir/pdwan-20140714.051601-data-A3-omp-1D.dat' u 1:3 t 'simple' w l lw 0.5 lc rgb 'blue', 'logDir/pdwan-20140714.051601-data-A3-omp-1D.dat' u 1:5 t 'dgemm' w l lw 0.5 lc rgb 'red' # pause -1 # Paula Dwan : Assignment 3 : A3-plotgraph-thread.gp reset set xtic auto set ytic auto set size 1,1 set grid set key outside # set title 'Comparison : No of Threads v Time taken' set ylabel 'Time taken / s' set xlabel 'no of Threads' set origin 0,0 set key outside plot 'logDir/pdwan-20140714.052102-data-A3-omp-1D.dat' u 2:3 t 'simple' w l lw 0.5 lc rgb 'blue', 'logDir/pdwan-20140714.052102-data-A3-omp-1D.dat' u 2:5 t 'dgemm' w l lw 0.5 lc rgb 'red' # pause -1 </pre>					

Plotting the following we get graphs for matrix size v time and also for no of threds v time (using -i so that each iteration uses the same values for A[i] and B[i] where I is the value of the column.

Matrix Size	no of Threads	Time / manual	Infinity Norm / manual	Time /dgemm	Infinity Norm / dgemm
10	2	0.000245	3025	0.000549	3025
10	2	0.002002	3025	0.002351	3025
10	2	0.001001	3025	0.001316	3025
20	2	0.001562	44100	0.002868	44100
20	2	0.000634	44100	0.001808	44100
20	2	0.000673	44100	0.001917	44100
30	10	0.000718	216225	0.002673	216225
30	10	0.000804	216225	0.003045	216225
30	10	0.000774	216225	0.003035	216225
40	10	0.001076	672400	0.004412	672400
40	10	0.001097	672400	0.004799	672400
40	10	0.001169	672400	0.004890	672400
50	10	0.001539	1.62562e+06	0.007291	1.62562e+06
50	10	0.001658	1.62562e+06	0.007325	1.62562e+06
50	10	0.001525	1.62562e+06	0.007136	1.62562e+06

Thankfully for Linux (Ubuntu) – I could install and run GNUplot locally.

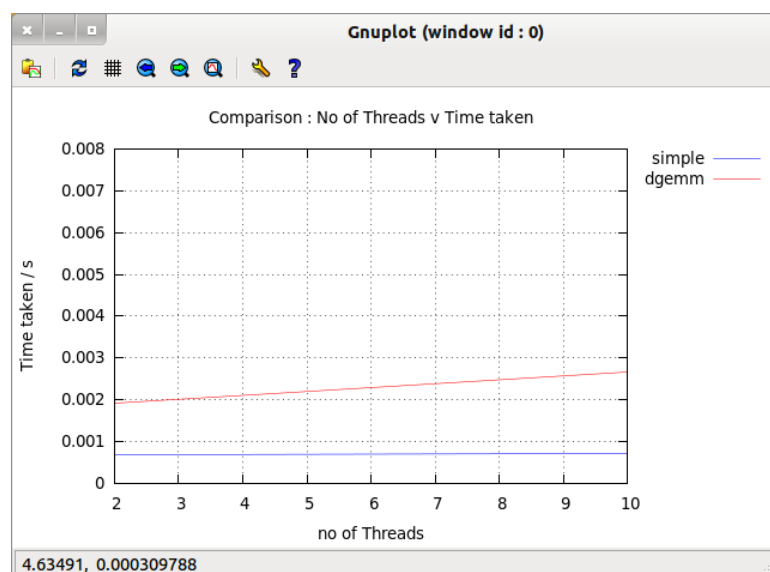
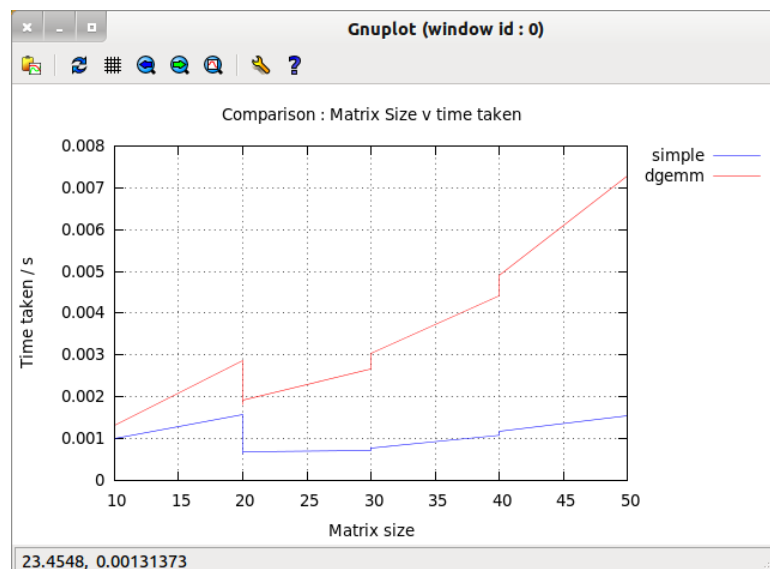
Screen shots of each were taken and added to the sections [GNUplot graphs.](#)

SUMMARY RESULTS :

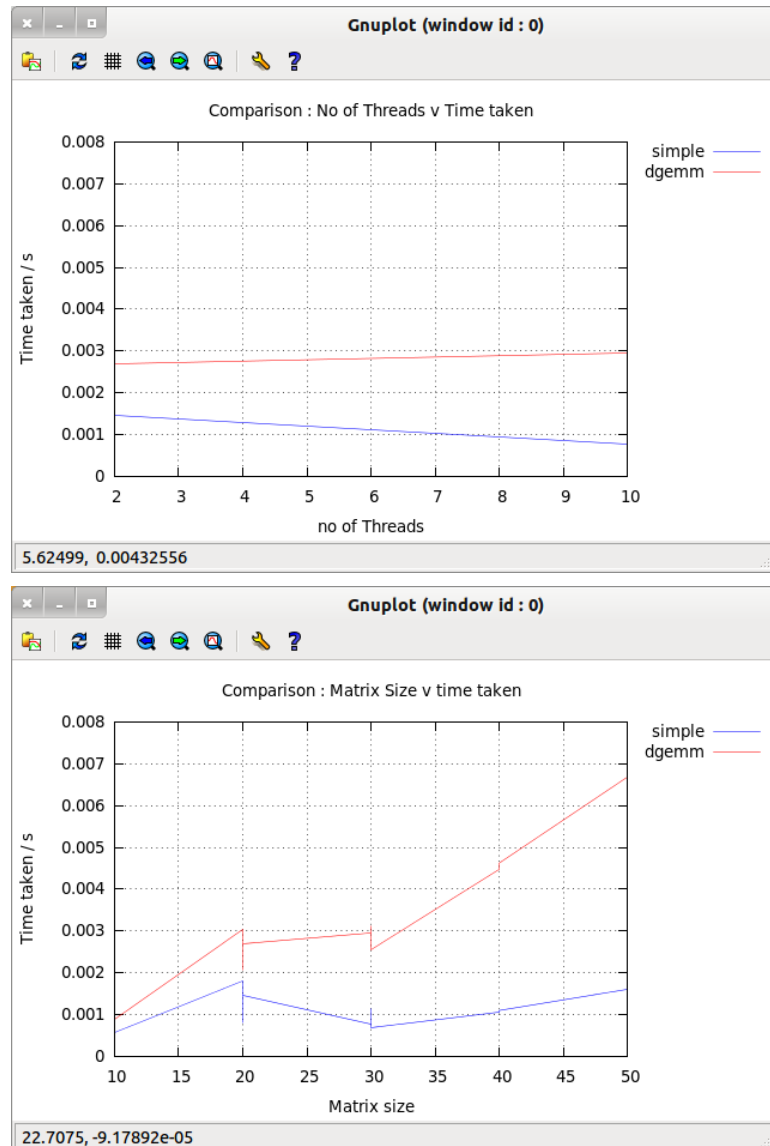
Build/plot:	<ul style="list-style-type: none"> The dependence of the execution time of the program on the matrix size n. The speedup over a serial counterpart of the program.
Variant :	<ul style="list-style-type: none"> One-step algorithm. No intermediate resulting matrix. Left matrix is horizontally partitioned The maximum absolute row sum norm (aka infinity-norm): $\ A\ _{\infty} = \max_{0 \leq i < n} \sum_{j=0}^{n-1} a_{ij} $
Infinity norm	<p>Sum the absolute values along each row and then take the biggest answer.</p> <p>Example: $A = \begin{vmatrix} 1 & -7 \\ -2 & -3 \end{vmatrix}$</p> <p>then matrix norm of A = $\max (1 + -7 , -2 + -3) = \max (8, 5) = \underline{\underline{8}}$</p>

GNU PLOT GRAPHS

Graph for logDir/pdwan-20140714.051601-data-A3-omp-1D.dat



Using different data set, the following is obtained :



CONCLUSIONS

The increase in size of the matrix is after a while no longer negates the increase in time taken. This is evident when the matrix size is increased to 1,000 x 1,000. Also the value was a little too large to graph. The following is an example of this.

Finally, there is only a finite number of threads available in any one system.

Matrix Size	no of Threads	Time / manual	Infinity Norm / manual	Time /dgemm	Infinity Norm / dgemm
50	10	0.001546	1.62562e+06	0.007661	1.62562e+06
50	10	0.001771	1.62562e+06	0.009032	1.62562e+06
50	10	0.001597	1.62562e+06	0.007557	1.62562e+06
100	10	0.007191	2.55025e+07	0.030856	2.55025e+07
100	10	0.010156	2.55025e+07	0.035289	2.55025e+07
100	10	0.008451	2.55025e+07	0.039047	2.55025e+07
500	20	1.363494	1.56876e+10	2.859575	1.56876e+10
500	20	1.370261	1.56876e+10	2.927285	1.56876e+10
500	20	1.325166	1.56876e+10	3.256134	1.56876e+10

Matrix Size	no of Threads	Time / manual	Infinity Norm / manual	Time /dgemm	Infinity Norm / dgemm
1000	20	12.946991	2.505e+11	21.963079	2.505e+11
1000	20	12.964027	2.505e+11	22.225876	2.505e+11
1000	20	12.597507	2.505e+11	22.171022	2.505e+11


```

pdwan@csserver:~/exercises/Assignment3
File Edit View Search Terminal Help
#
# Program : A3-omp-1D
# where : .dat contains timing data & .txt contains matrix values
#
# |Matrix| |Threads| Time/manual Inf Norm/manual Time/dgemm Inf Norm/dgemm
#
6 3 0.001521s 441 0.001662s 441
6 3 0.001246s 1396 0.001480s 1396
1,1 Top

```

Again, validation gives :

Source matrix initialized to value of row for each column										
A	0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9	10
1	1	2	3	4	5	6	7	8	9	10
2	1	2	3	4	5	6	7	8	9	10
3	1	2	3	4	5	6	7	8	9	10
4	1	2	3	4	5	6	7	8	9	10
5	1	2	3	4	5	6	7	8	9	10
6	1	2	3	4	5	6	7	8	9	10
7	1	2	3	4	5	6	7	8	9	10
8	1	2	3	4	5	6	7	8	9	10
9	1	2	3	4	5	6	7	8	9	10

Source matrix initialized to value of row for each column										
B	0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9	10
1	1	2	3	4	5	6	7	8	9	10
2	1	2	3	4	5	6	7	8	9	10
3	1	2	3	4	5	6	7	8	9	10
4	1	2	3	4	5	6	7	8	9	10
5	1	2	3	4	5	6	7	8	9	10
6	1	2	3	4	5	6	7	8	9	10
7	1	2	3	4	5	6	7	8	9	10
8	1	2	3	4	5	6	7	8	9	10
9	1	2	3	4	5	6	7	8	9	10

Program – calculate results using dot product										
Cijk	0	1	2	3	4	5	6	7	8	9
0	55	110	165	220	275	330	385	440	495	550
1	55	110	165	220	275	330	385	440	495	550
2	55	110	165	220	275	330	385	440	495	550
3	55	110	165	220	275	330	385	440	495	550
4	55	110	165	220	275	330	385	440	495	550
5	55	110	165	220	275	330	385	440	495	550
6	55	110	165	220	275	330	385	440	495	550
7	55	110	165	220	275	330	385	440	495	550
8	55	110	165	220	275	330	385	440	495	550
9	55	110	165	220	275	330	385	440	495	550

Infinity Norm : max of ltotal of each row\										
MANUAL										
	3025	3025	3025	3025	3025	3025	3025	3025	3025	3025
	3025	3025	3025	3025	3025	3025	3025	3025	3025	3025
	3025	3025	3025	3025	3025	3025	3025	3025	3025	3025
	3025	3025	3025	3025	3025	3025	3025	3025	3025	3025
	3025	3025	3025	3025	3025	3025	3025	3025	3025	3025
	3025	3025	3025	3025	3025	3025	3025	3025	3025	3025
	3025	3025	3025	3025	3025	3025	3025	3025	3025	3025
	3025	3025	3025	3025	3025	3025	3025	3025	3025	3025
	3025	3025	3025	3025	3025	3025	3025	3025	3025	3025

MANUAL										
Program – calculate results using for loops										
Cijk	0	1	2	3	4	5	6	7	8	9
0	55	110	165	220	275	330	385	440	495	550
1	55	110	165	220	275	330	385	440	495	550
2	55	110	165	220	275	330	385	440	495	550
3	55	110	165	220	275	330	385	440	495	550
4	55	110	165	220	275	330	385	440	495	550
5	55	110	165	220	275	330	385	440	495	550
6	55	110	165	220	275	330	385	440	495	550
7	55	110	165	220	275	330	385	440	495	550
8	55	110	165	220	275	330	385	440	495	550
9	55	110	165	220	275	330	385	440	495	550

CBLAS										
Program – calculate results using blas										
Cijk	0	1	2	3	4	5	6	7	8	9
0	55	110	165	220	275	330	385	440	495	550
1	55	110	165	220	275	330	385	440	495	550
2	55	110	165	220	275	330	385	440	495	550
3	55	110	165	220	275	330	385	440	495	550
4	55	110	165	220	275	330	385	440	495	550
5	55	110	165	220	275	330	385	440	495	550
6	55	110	165	220	275	330	385	440	495	550
7	55	110	165	220	275	330	385	440	495	550
8	55	110	165	220	275	330	385	440	495	550
9	55	110	165	220	275	330	385	440	495	550

Program – difference from dot-product										
Cijk	0	1	2	3	4	5	6	7	8	9
0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0
5	0	0	0	0	0	0	0	0	0	0
6	0	0	0	0	0	0	0	0	0	0
7	0	0	0	0	0	0	0	0	0	0
8	0	0	0	0	0	0	0	0	0	0
9	0	0	0	0	0	0	0	0	0	0

Program – difference from dot-product										
Cijk	0	1	2	3	4	5	6	7	8	9
0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0
5	0	0	0	0	0	0	0	0	0	0
6	0	0	0	0	0	0	0	0	0	0
7	0	0	0	0	0	0	0	0	0	0
8	0	0	0	0	0	0	0	0	0	0
9	0	0	0	0	0	0	0	0	0	0

APPENDIX II – ACKNOWLEDGEMENTS

- [wikipedia.org](https://www.wikipedia.org)
- Geln McLachlan – gnuplot tutorials (youtube)
- [Stackoverflow.com](https://stackoverflow.com) – examples of Open MP