

Cache Effect in Matrix Multiplication. Tutorial for Assignment 1


See also the original slides:

1. <http://www.eecs.harvard.edu/~mdw/course/cs61/mediawiki/images/b/bb/Lectures-cacheperf.pdf>
2. <http://www.cs.tau.ac.il/~ozery/courses/software-project09/presentations/assign2.ppt>

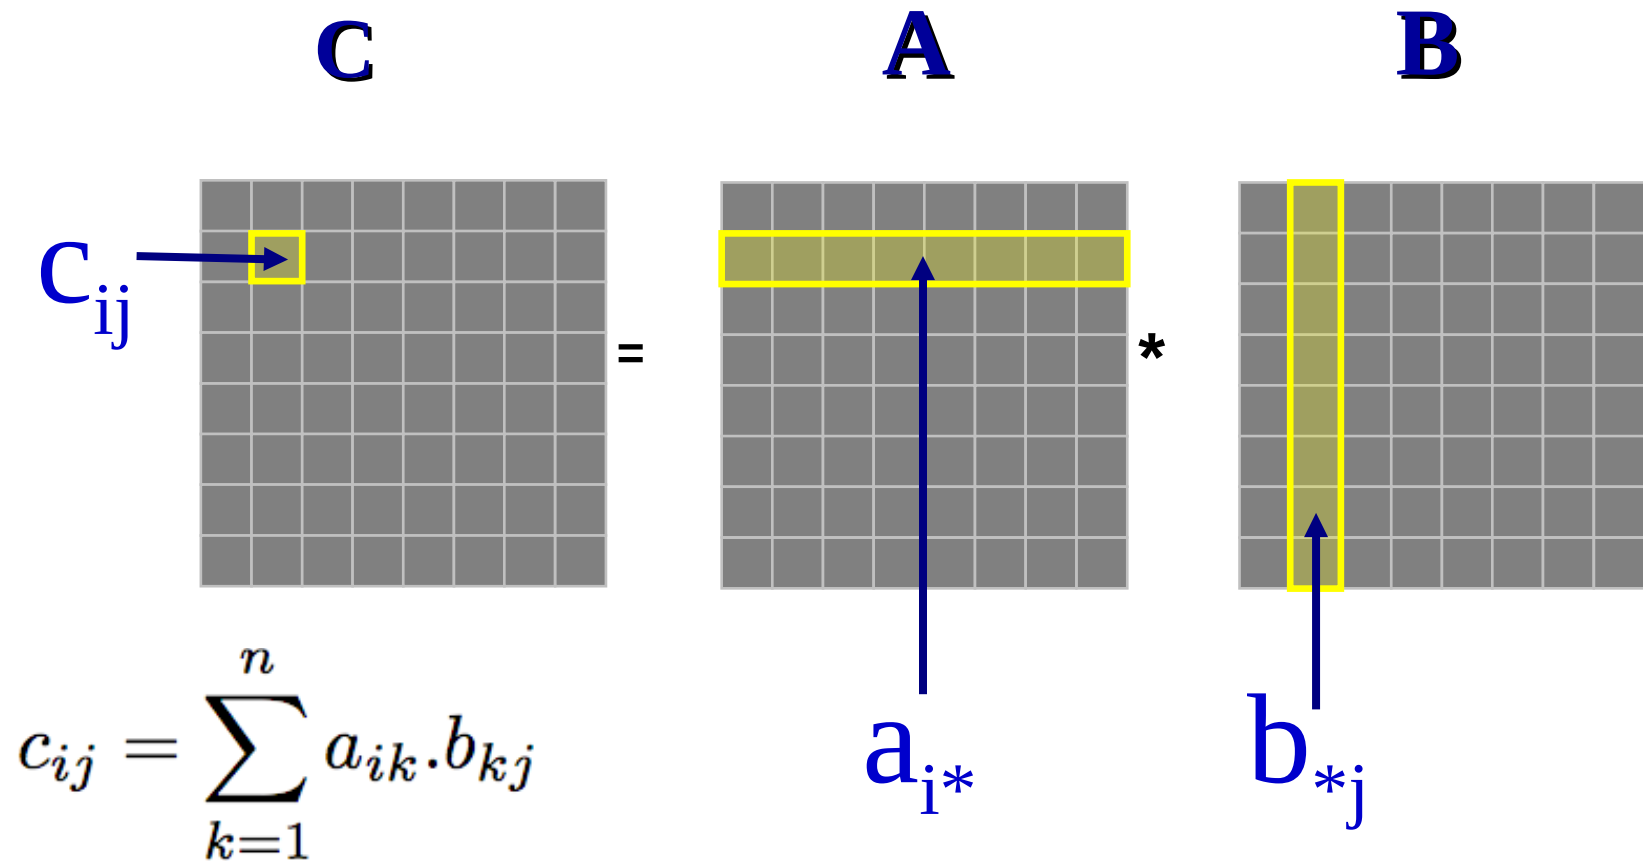
Multiplication of 2D Matrices

1. Simple algorithm
2. Time and Performance Measurement
3. Simple Code Improvements

Matrix Multiplication



Matrix Multiplication



The simplest algorithm

Assumption: the matrices are stored as 2-D $N \times N$ arrays

```
for (i=0; i<N; i++)  
    for (j=0; j<N; j++)  
        for (k=0; k<N; k++)  
            c[i][j] += a[i][k] * b[k][j];
```

Advantage: code simplicity

Disadvantage: performance (?)

First Improvement

```
for (i=0; i<N; i++)  
    for (j=0; j<N; j++)  
        for (k=0; k<N; k++)  
            c[i][j] += a[i][k] * b[k][j]  
];
```

c[i][j]:

- Requires address (pointer) computations
- is constant in the k-loop

First Performance Improvement

```
for (i=0;i<N;i++){  
    for (j=0;j<N;j++) {  
        int sum = 0;  
        for (k=0;k<N;k++) {  
            sum += a[i][k] * b[k][j];  
        }  
        c[i][j] = sum;  
    }  
}
```

Benchmarking

- ```
#include <sys/time.h>
struct timeval tv_start, tv_end;
struct timezone tz;
gettimeofday(&tv_start, &tz);
<code to benchmark>
gettimeofday(&tv_end, &tz);
double elapsed = (double) (tv_end.tv_sec -
tv_start.tv_sec) + (double) (tv_end.tv_usec -
tv_start.tv_usec) * 1.e-6;
```



# Cache Memory

# Cache memory (CPU cache)

- A temporary storage area where frequently accessed data can be stored for rapid access.
- Once the data is stored in the cache, future use can be made by accessing the cached copy rather than re-fetching the original data, so that the average access time is lower.

# Memory Hierarchy

## CPU Registers

500 Bytes  
0.25 ns

## Cache

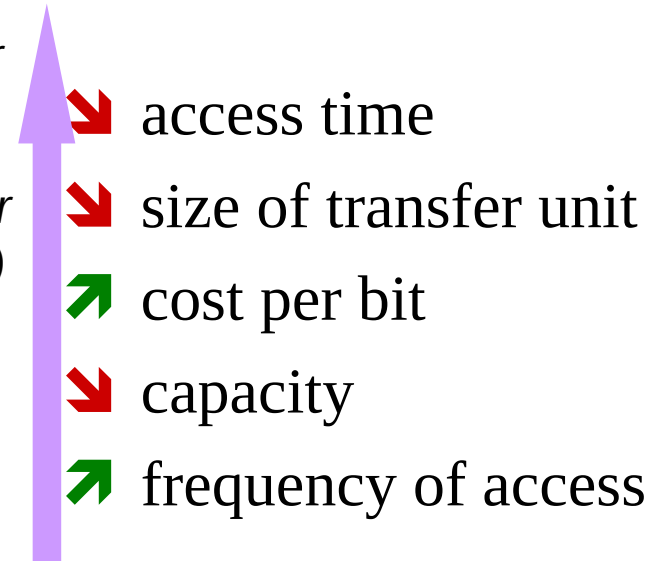

16K-1M Bytes  
1 ns

## Main Memory

64M-2G Bytes  
100ns

## Disk

100 G Bytes  
5 ms



- The memory cache is **closer** to the processor than the main memory.
- It is **smaller, faster and more expensive** than the main memory.
- Transfer between caches and main memory is performed in units called cache **blocks/lines**.

# Types of Cache Misses


1. **Compulsory misses:** caused by **first access** to blocks that have never been in cache (also known as *cold-start misses*)
2. **Capacity misses:** when cache cannot contain all the blocks needed during execution of a program. Occur because of blocks being replaced and later retrieved when accessed.
3. **Conflict misses:** when multiple blocks compete for the same set.

# Main Cache Principles

- **Temporal Locality** (Locality in Time): If an item is referenced, it will tend to be referenced again soon.
- **Spatial Locality** (Locality in Space): If an item is referenced, close items (by address) tend to be referenced soon.
- **Miss Rate** : Fraction of memory references not found in cache ( $\# \text{ misses} / \# \text{ references}$ )
- **Hit Rate**: Time to deliver a line in the cache to the processor (includes time to determine whether the line is in the cache)


# Improving Spatial Locality:

## Loop Reordering for Matrices Allocated by Row



# Writing Cache Friendly Code

- Assumptions: 1) block-size =  $\beta$  words (word = int)  
2)  $N$  is divided by  $\beta$   
3) The cache cannot hold a complete row / column



```
int sumarrayrows(int a[N][N])
{
 int i, j, sum = 0;
 for (i = 0; i < N; i++)
 for (j = 0; j < N; j++)
 sum += a[i][j];
 return sum;
}

int sumarraycols(int a[N][N])
{
 int i, j, sum = 0;
 for (j = 0; j < N; j++)
 for (i = 0; i < N; i++)
 sum += a[i][j];
 return sum;
}
```

Accesses successive elements:  
 $a[0][0], \dots, a[0][N-1], a[1][0], \dots$

15

Miss rate =  $1/\beta$

no spatial  
locality!

Accesses distant elements:  
 $a[0][0], \dots, a[N-1][0], a[0][1], \dots$

Miss rate = 1

# Matrix Multiplication (ijk)


```

/* ijk */
for (i=0; i<n; i++) {
 for (j=0; j<n; j++) {
 int sum = 0;
 for (k=0; k<n; k++)
 sum += a[i][k] * b[k][j];
 c[i][j] = sum;
 }
}

```

every element  
of A and B is  
accessed N  
times

Inner loop:



- Misses per Inner Loop Iteration ( $i=j=0$ ):

| <u>A</u>  | <u>B</u> | <u>C</u> |
|-----------|----------|----------|
| $1/\beta$ | 1.0      | 1.0      |

\*what happens when  $i=0$   
and  $j=1$ ?

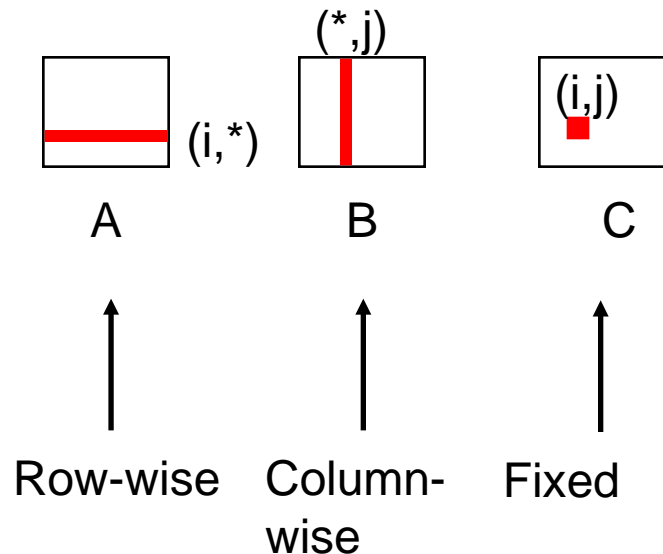


# Matrix Multiplication (jik)

every element  
of A and B is  
accessed N  
times

```
/* jik */
for (j=0; j<n; j++) {
 for (i=0; i<n; i++) {
 int sum = 0;
 for (k=0; k<n; k++)
 sum += a[i][k] * b[k][j];
 c[i][j] = sum;
 }
}
```

Inner loop:



- Misses per Inner Loop Iteration ( $i=j=0$ ):

| <u>A</u>  | <u>B</u> | <u>C</u> |
|-----------|----------|----------|
| $1/\beta$ | 1.0      | 1.0      |


\*what happens when  $j=0$   
and  $i=1$ ?

# Matrix Multiplication (kij)

every element  
of B and C is  
accessed N  
times

```
/* kij */
for (k=0; k<n; k++) {
 for (i=0; i<n; i++) {
 int x = a[i][k];
 for (j=0; j<n; j++)
 c[i][j] += x * b[k][j];
 }
}
```

Inner loop:



- Misses per Inner Loop Iteration ( $i=j=0$ ):

| <u>A</u> | <u>B</u>  | <u>C</u>  |
|----------|-----------|-----------|
| 1.0      | $1/\beta$ | $1/\beta$ |


\*what happens when  $k=0$   
and  $i=1$ ?

# Matrix Multiplication (jki)

every element  
of A and C is  
accessed N  
times


```
/* jki */
for (j=0; j<n; j++) {
 for (k=0; k<n; k++) {
 int x = b[k][j];
 for (i=0; i<n; i++)
 c[i][j] += a[i][k] * x;
 }
}
```

Inner loop:




A

Column -  
wise



B

Fixed



C

Column-  
wise

- Misses per Inner Loop Iteration (i=j=0):

| <u>A</u> | <u>B</u> | <u>C</u> |
|----------|----------|----------|
| 1.0      | 1.0      | 1.0      |

\*what happens when j=0  
and k=1?

# Summary: misses ratios for the first iteration of the inner loop

**ijk (& jik):  $(\beta+1)/(2\beta) > 1/2$     kij:  $1/\beta$**

**jki: 1**


```
for (i=0; i<n; i++) {
 for (j=0; j<n; j++) {
 int sum = 0;
 for (k=0; k<n; k++)
 sum += a[i][k] * b[k][j];
 c[i][j] = sum;
 }
}
```

```
for (k=0; k<n; k++) {
 for (i=0; i<n; i++) {
 int x = a[i][k];
 for (j=0; j<n; j++)
 c[i][j] += x * b[k][j];
 }
}
```

```
for (j=0; j<n; j++) {
 for (k=0; k<n; k++) {
 int x = b[k][j];
 for (i=0; i<n; i++)
 c[i][j] += a[i][k] * x;
 }
}
```

# Improving Temporal Locality: Blocked Matrix Multiplication


# “Blocked” Matrix Multiplication



$$c_{ij} = \sum_{k=1}^N a_{ik} b_{kj}$$

**Key idea:** reuse the other elements in each cache block as much as possible


# “Blocked” Matrix Multiplication



The blocks loaded for the computation of  $C[i][j]$  are appropriate for the computation of  $C[i,j+1] \dots C[i,j+b-1]$

- compute the first  $b$  terms of  $C[i][j], \dots, C[i][j+b-1]$
- compute the next  $b$  terms of  $C[i][j], \dots, C[i][j+b-1]$

# “Blocked” Matrix Multiplication




**Next improvement:**

Reuse the loaded blocks of B for the computation of next  $(b - 1)$  subrows.



# “Blocked” Matrix Multiplication



Order of the operations:

Compute the first b terms of  $C_1$  ( $=A_1*B_1$ )

Compute the next b terms of  $C_1$  ( $=A_2*B_2$ )

...

Compute the last b terms of  $C_1$  ( $=A_4*B_4$ )

# “Blocked” Matrix Multiplication

|          |          |          |          |
|----------|----------|----------|----------|
| $C_{11}$ | $C_{12}$ | $C_{13}$ | $C_{14}$ |
| $C_{21}$ | $C_{22}$ | $C_{23}$ | $C_{24}$ |
| $C_{31}$ | $C_{32}$ | $C_{33}$ | $C_{34}$ |
| $C_{41}$ | $C_{42}$ | $C_{43}$ | $C_{44}$ |

|          |          |          |          |
|----------|----------|----------|----------|
| $A_{11}$ | $A_{12}$ | $A_{13}$ | $A_{14}$ |
| $A_{21}$ | $A_{22}$ | $A_{23}$ | $A_{24}$ |
| $A_{31}$ | $A_{32}$ | $A_{33}$ | $A_{34}$ |
| $A_{41}$ | $A_{42}$ | $A_{43}$ | $A_{44}$ |

|          |          |          |          |
|----------|----------|----------|----------|
| $B_{11}$ | $B_{12}$ | $B_{13}$ | $B_{14}$ |
| $B_{21}$ | $B_{22}$ | $B_{23}$ | $B_{24}$ |
| $B_{31}$ | $B_{32}$ | $B_{33}$ | $B_{34}$ |
| $B_{41}$ | $B_{42}$ | $B_{43}$ | $B_{44}$ |

$$C_{22} = A_{21}B_{12} + A_{22}B_{22} + A_{23}B_{32} + A_{24}B_{42} = \sum_k A_{2k} * B_{k2}$$

$$N = 4 * b$$

**Main Point:** each multiplication operates on small “block” matrices, whose size may be chosen so that they fit in the cache.

# Blocked Algorithm

- The blocked version of the i-j-k algorithm is written simply as

```
for (i=0;i<N/b;i++)
 for (j=0;j<N/b;j++)
 for (k=0;k<N/b;k++)
 C[i][j] += A[i][k]*B[k][j]
```

b x b matrix addition

b x b matrix multiplication

- $b$  = block (sub-matrix) size (Assume  $b$  divides  $N$ )
- $X[i][j]$  = a sub-matrix of  $X$ , defined by block row  $i$  and block column  $j$

# Maximum Block Size

- The blocking optimization works only if the **blocks fit in cache**.
- That is, **3** blocks of size  **$b \times b$**  must fit in memory (for A, B, and C)
- **$M$**  = size of cache (in elements/words)
- We must have:  **$3b^2 \approx M$ , or  $b \approx \sqrt{M/3}$**
- **Lower bound** =  $(b^2/\beta) (2(n/b)^3 + (n/b)^2) = (1/\beta)(2n^3/b + n^2) = \Theta(n^3/(b\beta)) = \Theta(n^3/(\beta\sqrt{M}))$
- Therefore, the ratio of cache misses ijk-blocked vs. ijk-unblocked:  **$1:\sqrt{M}$**