

We consider a Pthreads application computing the dot product of two real  $m$ -length vectors  $x$  and  $y$  to illustrate parallel programming an  $n$ -processor SMP computer with thread libraries. This MT application divides vectors  $x$  and  $y$  into  $n$  sub-vectors. Meanwhile the first  $n-1$  sub-vectors are of the same length  $\frac{m}{n}$ , the last  $n$ -th subvector may be shorter if  $m$  is not a multiply of  $n$ . This application uses  $n$  parallel threads with  $i$ -th thread computing its fraction of the total dot product by multiplying sub-vectors  $x_i$  and  $y_i$ . The  $n$  parallel threads share a data object accumulating the dot product, and synchronize their access to the data object with a mutex. The main thread creates the  $n$  threads, waits for them to complete their computations via joining with each of the threads, and then outputs the result.

The source code of the application is as follows:

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

#define MAXTHRDS 124

typedef struct {
    double *my_x;
    double *my_y;
    double my_dot_prod;
    double *global_dot_prod;
    pthread_mutex_t *mutex;
    int my_vec_len;
} dot_product_t;

void *serial_dot_product(void *arg) {
    dot_product_t *dot_data;
    int i;

    dot_data = arg;
    for(i=0; i<dot_data->my_vec_len; i++)
        dot_data->my_dot_prod += dot_data->my_x[i]*dot_data->my_y[i];

    pthread_mutex_lock(dot_data->mutex);
    *(dot_data->global_dot_prod) += dot_data->my_dot_prod;
    pthread_mutex_unlock(dot_data->mutex);

    pthread_exit(NULL);
}

int main()
{
    double *x, *y, dot_prod;
    pthread_t *working_thread;
    dot_product_t *thrd_dot_prod_data;
    void *status;
    pthread_mutex_t *mutex_dot_prod;

    int num_of_thrds;
    int vec_len;
```

```

int subvec_len;
int i;

printf("Number of processors = ");
if(scanf("%d", &num_of_thrds) < 1 || num_of_thrds > MAXTHRDS) {
    printf("Check input for number of processors. Bye.\n");
    return -1;
}
printf("Vector length = ");
if(scanf("%d", &vec_len)<1) {
    printf("Check input for vector length. Bye.\n");
    return -1;
}
subvec_len = vec_len/num_of_thrds;

x = malloc(vec_len*sizeof(double));
y = malloc(vec_len*sizeof(double));
for(i=0; i<vec_len; i++) {
    x[i] = 1.;
    y[i] = 1.;
}

working_thread = malloc(num_of_thrds*sizeof(pthread_t));
thrd_dot_prod_data = malloc(num_of_thrds*sizeof(dot_product_t));

mutex_dot_prod = malloc(sizeof(pthread_mutex_t));
pthread_mutex_init(mutex_dot_prod, NULL);

for(i=0; i<num_of_thrds; i++) {
    thrd_dot_prod_data[i].my_x = x + i*subvec_len;
    thrd_dot_prod_data[i].my_y = y + i*subvec_len;
    thrd_dot_prod_data[i].my_dot_prod = 0.;
    thrd_dot_prod_data[i].global_dot_prod = &dot_prod;
    thrd_dot_prod_data[i].mutex = mutex_dot_prod;
    thrd_dot_prod_data[i].my_vec_len =
        (i==num_of_thrds-1)? vec_len-(num_of_thrds-1)*subvec_len
        : subvec_len;
    pthread_create(&working_thread[i], NULL, serial_dot_product,
        (void*)&thrd_dot_prod_data[i]);
}
for(i=0; i<num_of_thrds; i++)
    pthread_join(working_thread[i], &status);

printf("Dot product = %f\n", dot_prod);

free(x);
free(y);
free(working_thread);
free(thrd_dot_prod_data);
pthread_mutex_destroy(mutex_dot_prod);
free(mutex_dot_prod);
}

```