

```

#include <stdio.h>
#include <stdlib.h>
#include <float.h>
#include <mpi.h>

int n, p;

int main(int argc, char **argv) {
    int myn, myrank;
    double *a, *b, *c, *allB, start, sum, *allC, sumdiag;
    int i, j, k;

    n = atoi(argv[1]);
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &p);
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
    myn = n/p;
    a = malloc(myn*n*sizeof(double));
    b = malloc(myn*n*sizeof(double));
    c = malloc(myn*n*sizeof(double));
    allB = malloc(n*n*sizeof(double));
    for(i=0; i<myn*n; i++) {
        a[i] = 1.;
        b[i] = 2.;
    }
    MPI_Barrier(MPI_COMM_WORLD);
    if(myrank==0)
        start = MPI_Wtime();
    for(i=0; i<p; i++)
        MPI_Gather(b, myn*n, MPI_DOUBLE, allB, myn*n, MPI_DOUBLE,
            i, MPI_COMM_WORLD);
    for(i=0; i<myn; i++)
        for(j=0; j<n; j++) {
            sum = 0.;
            for(k=0; k<n; k++)
                sum += a[i*n+k]*allB[k*n+j];
            c[i*n+j] = sum;
        }
    free(allB);
    MPI_Barrier(MPI_COMM_WORLD);
    if(myrank==0)
        printf("It took %f seconds to multiply 2 %dx%d matrices.\n",
            MPI_Wtime()-start, n, n);
    if(myrank==0)
        allC = malloc(n*n*sizeof(double));
    MPI_Gather(c, myn*n, MPI_DOUBLE, allC, myn*n, MPI_DOUBLE,
        0, MPI_COMM_WORLD);
    if(myrank==0) {
        for(i=0, sumdiag=0.; i<n; i++)
            sumdiag += allC[i*n+i];
        printf("The trace of the resulting matrix is %f\n", sumdiag);
    }
    if(myrank==0)
        free(allC);
    MPI_Finalize();
}

```

```

    free(a);
    free(b);
    free(c);
}

```

It is assumed that the user via an external parameter supplies the value for  $n$  and this value is a multiple of the total number of processes executing the program,  $p$ . It is also assumed that process zero ( $\text{myrank}=0$ ) runs on the node with I/O available. Each process has the following arrays on heap:

- $a$ ,  $b$ , and  $c$  to store its horizontal  $\text{myn} \times n$  slice of matrices  $A$ ,  $B$ , and  $C$  respectively;
- $\text{allB}$  to store the full matrix  $B$ . This array is deallocated as soon as the process has computed all elements of its  $C$  slice.

The array  $\text{allC}$  is allocated on process zero after the parallel multiplication itself has been complete. This array is used as a receive buffer by the collective operation that gathers all slices of the resulting matrix  $C$  on process zero. After the matrix  $C$  is gathered on process zero in the array  $\text{allC}$ , this process computes the trace of the matrix and output it (just to let the user check correctness of the computations).

By the way, the loop

```

for(i=0; i<p; i++)
    MPI_Gather(b, myn*n, MPI_DOUBLE, allB, myn*n, MPI_DOUBLE,
              i, MPI_COMM_WORLD);

```

resulting in all processes to receive the matrix  $B$ , can be replaced by a single call of the form

```

MPI_Allgather(b, myn*n, MPI_DOUBLE, allB, myn*n, MPI_DOUBLE,
             MPI_COMM_WORLD);

```

$\text{MPI\_Allgather}$  is an example of advanced variation of the basic gather operation. It is assumed that MPI may implement  $\text{MPI\_Allgather}$  more efficiently than via multiple calls to  $\text{MPI\_Gather}$ .